

dsa

October 21, 2019

0.1 IOT Data Cleaning and Visualisation

0.1.1 Import required packages

```
In [34]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.dates as mdates
import re
import os
import datetime
import seaborn as sn
```

0.1.2 Reading in data obtained from IOT Platform

- Download data as csv files and use load it using pandas
- Download json data through the API
- download the sample air quality data csv file from the session website (session 4). (<https://airqo.net/dsa/>)

```
In [35]: base_dir = 'data'
airquality_data = pd.read_csv(os.path.join(base_dir, 'AQ_15_DSA.csv'), parse_dates = [
```

0.1.3 Data Inspection

```
In [36]: airquality_data.head()
```

```
Out[36]:
```

	created_at	entry_id	field1	field2	field3	field4	field5	\
0	2019-01-29 12:40:18	1	25.48	27.33	28.52	30.42	0.000000	
1	2019-01-29 12:41:34	2	24.87	25.83	31.00	32.65	1000.000000	
2	2019-01-29 12:47:33	3	28.52	31.03	29.80	31.35	0.000000	
3	2019-01-29 12:49:56	4	27.15	28.35	30.80	32.10	1000.000000	
4	2019-01-31 10:40:45	5	17.68	19.68	17.47	19.07	1.333921	

	field6	field7	field8	latitude	longitude	elevation	status
0	0.000000	4.02	NaN	NaN	NaN	NaN	NaN
1	1000.000000	3.79	NaN	NaN	NaN	NaN	NaN
2	0.000000	4.02	NaN	NaN	NaN	NaN	NaN

3	1000.000000	3.77	NaN	NaN	NaN	NaN	NaN
4	34.372246	3.88	NaN	NaN	NaN	NaN	NaN

In [37]: `airquality_data.tail()`

```
Out[37]:
```

		created_at	entry_id	field1	field2	field3	field4	field5	\
93829	2019-10-02	14:40:36	93830	29.60	40.52	31.43	36.25	1.333871	
93830	2019-10-02	14:41:59	93831	101.13	116.60	105.43	112.90	1.333883	
93831	2019-10-02	14:43:29	93832	142.78	156.87	155.25	165.80	1.333903	
93832	2019-10-02	14:44:49	93833	35.65	42.58	34.17	39.97	1.333855	
93833	2019-10-02	14:46:11	93834	40.25	52.53	41.60	50.95	1.333859	

	field6	field7	field8	\
93829	34.372238	3.71	1.333871,34.372238,1772.30,0.15,10.00,88.00,0...	
93830	34.372231	3.71	1.333883,34.372231,1774.10,0.05,10.00,84.00,0...	
93831	34.372280	3.70	1.333903,34.372280,1774.90,0.07,10.00,100.00,0...	
93832	34.372295	3.70	1.333855,34.372295,1764.10,0.08,11.00,83.00,0...	
93833	34.372288	3.70	1.333859,34.372288,1760.50,0.03,11.00,83.00,0...	

	latitude	longitude	elevation	status
93829	NaN	NaN	NaN	NaN
93830	NaN	NaN	NaN	NaN
93831	NaN	NaN	NaN	NaN
93832	NaN	NaN	NaN	NaN
93833	NaN	NaN	NaN	NaN

```
In [38]: print('shape of the data')
print(airquality_data.shape)

print('number of rows in the dataset')
print(airquality_data.shape[0])
```

```
shape of the data
(93834, 14)
number of rows in the dataset
93834
```

```
In [39]: print(list(airquality_data.columns.values))
```

```
['created_at', 'entry_id', 'field1', 'field2', 'field3', 'field4', 'field5', 'field6', 'field7']
```

```
In [40]: airquality_data.dtypes
```

```
Out[40]: created_at    datetime64[ns]
entry_id              int64
field1                float64
field2                float64
```

```

field3          float64
field4          float64
field5          float64
field6          float64
field7          float64
field8          object
latitude        float64
longitude        float64
elevation        float64
status          float64
dtype: object

```

```
In [41]: airquality_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 93834 entries, 0 to 93833
Data columns (total 14 columns):
created_at      93834 non-null datetime64[ns]
entry_id        93834 non-null int64
field1          93834 non-null float64
field2          93834 non-null float64
field3          93834 non-null float64
field4          93834 non-null float64
field5          93834 non-null float64
field6          93834 non-null float64
field7          93834 non-null float64
field8          82085 non-null object
latitude        0 non-null float64
longitude       0 non-null float64
elevation       0 non-null float64
status          0 non-null float64
dtypes: datetime64[ns](1), float64(11), int64(1), object(1)
memory usage: 10.0+ MB

```

0.1.4 Rename column names

```
In [42]: airquality_data.rename(columns={'created_at': 'TimeStamp', 'field1': 's1_pm2_5', 'field2':
                                         'field3': 's2_pm2_5', 'field4': 's2_pm10',
                                         'field5': 'Latitude', 'field6': 'Longitude', 'fie
```

0.1.5 Drop the columns that we won't be using

```
In [43]: airquality_data_of_interest = airquality_data.drop(['entry_id', 'Battery Voltage', 'lati
                                                             'status', 'GpsData'], axis=1)
```

0.1.6 Show the summaries i.e mean, median, quaurtiles, count

```
In [44]: airquality_data_of_interest.describe()
```

```
Out [44]:
```

	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude \
count	93834.000000	93834.000000	93834.000000	93834.000000	93834.000000
mean	21.664234	27.359923	23.249641	26.12310	597.178577
std	15.953107	20.357442	16.550017	19.46318	489.928987
min	0.000000	0.000000	0.000000	0.00000	0.000000
25%	12.220000	14.930000	13.170000	14.28000	1.333874
50%	18.280000	22.830000	19.950000	21.53000	1000.000000
75%	27.400000	34.580000	29.570000	32.50000	1000.000000
max	891.520000	1490.320000	725.800000	1082.08000	1000.000000

	Longitude
count	93834.000000
mean	610.290555
std	473.989563
min	0.000000
25%	34.372250
50%	1000.000000
75%	1000.000000
max	1000.000000

```
In [45]: airquality_data_of_interest.head()
```

```
Out [45]:
```

	TimeStamp	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude \
0	2019-01-29 12:40:18	25.48	27.33	28.52	30.42	0.000000
1	2019-01-29 12:41:34	24.87	25.83	31.00	32.65	1000.000000
2	2019-01-29 12:47:33	28.52	31.03	29.80	31.35	0.000000
3	2019-01-29 12:49:56	27.15	28.35	30.80	32.10	1000.000000
4	2019-01-31 10:40:45	17.68	19.68	17.47	19.07	1.333921

	Longitude
0	0.000000
1	1000.000000
2	0.000000
3	1000.000000
4	34.372246

0.1.7 Parsing the created_at field into a datetime object

This is important for indexing using datetime

localise timezone

In case of any duplicate timestamp, keep one of the duplicated records

```
In [46]: airquality_data_of_interest['TimeStamp'].dt.tz_localize('Africa/Kampala')
airquality_data_of_interest = airquality_data_of_interest.drop_duplicates(subset='TimeStamp')
time_indexed_data = airquality_data_of_interest.set_index('TimeStamp')
time_indexed_data.head()
```

```
Out [46]:
```

	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude \
TimeStamp					

2019-01-29 12:40:18	25.48	27.33	28.52	30.42	0.000000
2019-01-29 12:41:34	24.87	25.83	31.00	32.65	1000.000000
2019-01-29 12:47:33	28.52	31.03	29.80	31.35	0.000000
2019-01-29 12:49:56	27.15	28.35	30.80	32.10	1000.000000
2019-01-31 10:40:45	17.68	19.68	17.47	19.07	1.333921

	Longitude
TimeStamp	
2019-01-29 12:40:18	0.000000
2019-01-29 12:41:34	1000.000000
2019-01-29 12:47:33	0.000000
2019-01-29 12:49:56	1000.000000
2019-01-31 10:40:45	34.372246

0.1.8 Get average values for both sensors

0.1.9 Deal with outliers

In [47]: *# Finding out how many latitude or longitude entries have 0 or 1000*

```
bad_lat = []
bad_long = []

for lat in time_indexed_data['Latitude']:
    if ((lat==0) or (lat==1000)):
        bad_lat.append(lat)
print ('Bad latitudes are ', len(bad_lat))
for long in time_indexed_data['Longitude']:
    if ((long==0) or (long==1000)):
        bad_long.append(long)
print ('Bad longitudes are ', len(bad_long))
```

Bad latitudes are 56594
Bad longitudes are 56594

In [48]: *#Replacing latitude and longitude values of 0 and 1000 with Null*

```
time_indexed_data['Latitude'] = time_indexed_data['Latitude'].replace([0,1000], np.nan)
time_indexed_data['Longitude'] =time_indexed_data['Longitude'].replace([0,1000], np.nan)
time_indexed_data.head(5)
```

Out[48]:

	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude	Longitude
TimeStamp						
2019-01-29 12:40:18	25.48	27.33	28.52	30.42	NaN	NaN
2019-01-29 12:41:34	24.87	25.83	31.00	32.65	NaN	NaN
2019-01-29 12:47:33	28.52	31.03	29.80	31.35	NaN	NaN
2019-01-29 12:49:56	27.15	28.35	30.80	32.10	NaN	NaN
2019-01-31 10:40:45	17.68	19.68	17.47	19.07	1.333921	34.372246

0.1.10 Show null values in each column

```
In [49]: time_indexed_data.isnull().sum()
```

```
Out[49]: s1_pm2_5      0
         s1_pm10      0
         s2_pm2_5      0
         s2_pm10      0
         Latitude    56594
         Longitude    56594
         dtype: int64
```

```
In [50]: time_indexed_data['Latitude'] = time_indexed_data['Latitude'].fillna(method = 'bfill')
         time_indexed_data['Longitude'] = time_indexed_data['Longitude'].fillna(method = 'bfill')
         time_indexed_data.isnull().sum()
```

```
Out[50]: s1_pm2_5      0
         s1_pm10      0
         s2_pm2_5      0
         s2_pm10      0
         Latitude      0
         Longitude      0
         dtype: int64
```

```
In [51]: #Replacing values outside acceptable range with the other sensor's data
```

```
time_indexed_data['s1_pm2_5'] = np.where(((time_indexed_data['s1_pm2_5'] <= 0) | (time_indexed_data['s1_pm2_5'] > 500.4)) & (time_indexed_data['s2_pm2_5'] > 0), time_indexed_data['s2_pm2_5'], time_indexed_data['s1_pm2_5'])
```

```
time_indexed_data['s2_pm2_5'] = np.where(((time_indexed_data['s2_pm2_5'] <= 0) | (time_indexed_data['s2_pm2_5'] > 500.4)) & (time_indexed_data['s1_pm2_5'] > 0), time_indexed_data['s1_pm2_5'], time_indexed_data['s2_pm2_5'])
```

```
In [52]: #Dropping pm2.5 greater than 500.4 or less than or equal to 0 for both sensors
```

```
outlier_indices = time_indexed_data[((time_indexed_data['s1_pm2_5'] > 500.4) & (time_indexed_data['s1_pm2_5'] <= 0)) | ((time_indexed_data['s2_pm2_5'] > 500.4) & (time_indexed_data['s2_pm2_5'] <= 0))]
print(len(outlier_indices))
time_indexed_data.drop(outlier_indices, inplace = True)
time_indexed_data.info()
```

```
4
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
DatetimeIndex: 93830 entries, 2019-01-29 12:40:18 to 2019-10-02 14:46:11
```

```
Data columns (total 6 columns):
```

```
s1_pm2_5      93830 non-null float64
s1_pm10      93830 non-null float64
s2_pm2_5      93830 non-null float64
s2_pm10      93830 non-null float64
Latitude      93830 non-null float64
```

```
Longitude      93830 non-null float64
dtypes: float64(6)
memory usage: 5.0 MB
```

```
In [53]: time_indexed_data.head()
```

```
Out [53]:
```


	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude	Longitude
TimeStamp						
2019-01-29 12:40:18	25.48	27.33	28.52	30.42	1.333921	34.372246
2019-01-29 12:41:34	24.87	25.83	31.00	32.65	1.333921	34.372246
2019-01-29 12:47:33	28.52	31.03	29.80	31.35	1.333921	34.372246
2019-01-29 12:49:56	27.15	28.35	30.80	32.10	1.333921	34.372246
2019-01-31 10:40:45	17.68	19.68	17.47	19.07	1.333921	34.372246

```
In [54]: ## obtaining the averages of sensor one and sensor two readings
time_indexed_data['pm2_5'] =time_indexed_data[['s1_pm2_5', 's2_pm2_5']].mean(axis=1)
time_indexed_data['pm10'] =time_indexed_data[['s1_pm10', 's2_pm10']].mean(axis=1)
```

0.1.11 Data Visualisation

```
In [55]: fig = plt.figure(figsize=(14,6))
plt.rcParams.update({'font.size':14, 'font.weight':'bold'})
plt.plot(time_indexed_data.index,
         time_indexed_data['pm2_5'], color='purple', linestyle='solid')

chart_title = 'PM2.5 Concentration For Location X'
plt.title(chart_title,fontsize=20)
plt.ylabel('PM2.5', fontsize=20)
plt.xlabel('Date', fontsize=20)
plt.xticks(rotation=0)
plt.legend()
plt.show()
plt.tight_layout()
fig.savefig(os.path.join(base_dir,'pm25_concentration_for_location_x.png'))
```



<Figure size 432x288 with 0 Axes>

0.1.12 Compute average, minimum and maximum hourly, daily, monthly concentration values

- calculate the corresponding averages for hourly, daily and monthly observations

```
In [56]: hourly_average_airquality_data_concentrations = time_indexed_data.resample('H').mean()
hourly_maximum_airquality_data_concentrations = time_indexed_data.resample('H').max()
hourly_minimum_airquality_data_concentrations = time_indexed_data.resample('H').min()

daily_average_airquality_data_concentrations = time_indexed_data.resample('D').mean()
daily_maximum_airquality_data_concentrations = time_indexed_data.resample('D').max()
daily_minimum_airquality_data_concentrations = time_indexed_data.resample('D').min()

monthly_average_airquality_data_concentrations = time_indexed_data.resample('M').mean()
monthly_maximum_airquality_data_concentrations = time_indexed_data.resample('M').max()
monthly_minimum_airquality_data_concentrations = time_indexed_data.resample('M').min()

hourly_average_airquality_data_concentrations.head()
```

```
Out [56]:
```

	s1_pm2_5	s1_pm10	s2_pm2_5	s2_pm10	Latitude	\
TimeStamp						
2019-01-29 12:00:00	26.51	28.14	30.03	31.63	1.33	
2019-01-31 10:00:00	18.47	21.23	18.91	21.35	1.33	
2019-01-31 11:00:00	7.02	7.95	8.14	8.74	1.33	
2019-01-31 12:00:00	9.70	10.25	11.01	11.47	1.33	
2019-01-31 13:00:00	7.81	8.56	8.65	9.14	1.33	

```

Longitude  pm2_5  pm10

```


TimeStamp				
2019-01-29 12:00:00	34.37	28.27	29.88	
2019-01-31 10:00:00	34.37	18.69	21.29	
2019-01-31 11:00:00	34.37	7.58	8.34	
2019-01-31 12:00:00	34.37	10.35	10.86	
2019-01-31 13:00:00	34.37	8.23	8.85	

```
In [57]: plt.rcParams.update({'font.size':14, 'font.weight':'bold','mathtext.default': 'r'
x_axis_label = 'Months'
y_axis_label = 'PM25 and PM10 Concentration'
chart_title = 'Monthly Average $PM_{2.5}$ and $PM_{10}$ for location X'
file_path_no_aqi_color = os.path.join(base_dir , 'average_pm25_foreach_month_


fig = plt.figure(figsize=(12,8))
plt.rcParams.update({'font.size':12, 'font.weight':'bold', 'mathtext.default': 'r'
ax = fig.add_subplot(111)

plt.plot(monthly_average_airquality_data_concentrations.index,
         monthly_average_airquality_data_concentrations['pm2_5'].values,
         color='black', marker='P', linestyle='solid', label='PM2.5', linewidth=1.5)

plt.plot(monthly_average_airquality_data_concentrations.index,
         monthly_average_airquality_data_concentrations['pm10'].values,
         color='blue', marker='o', linestyle='solid', label='PM10', linewidth=1.5)

plt.title(chart_title, fontsize=20)
plt.ylabel(y_axis_label, fontsize=20)
plt.xlabel(x_axis_label,fontsize=20)

plt.legend(loc='upper right')
plt.show()
fig.savefig(file_path_no_aqi_color)
```

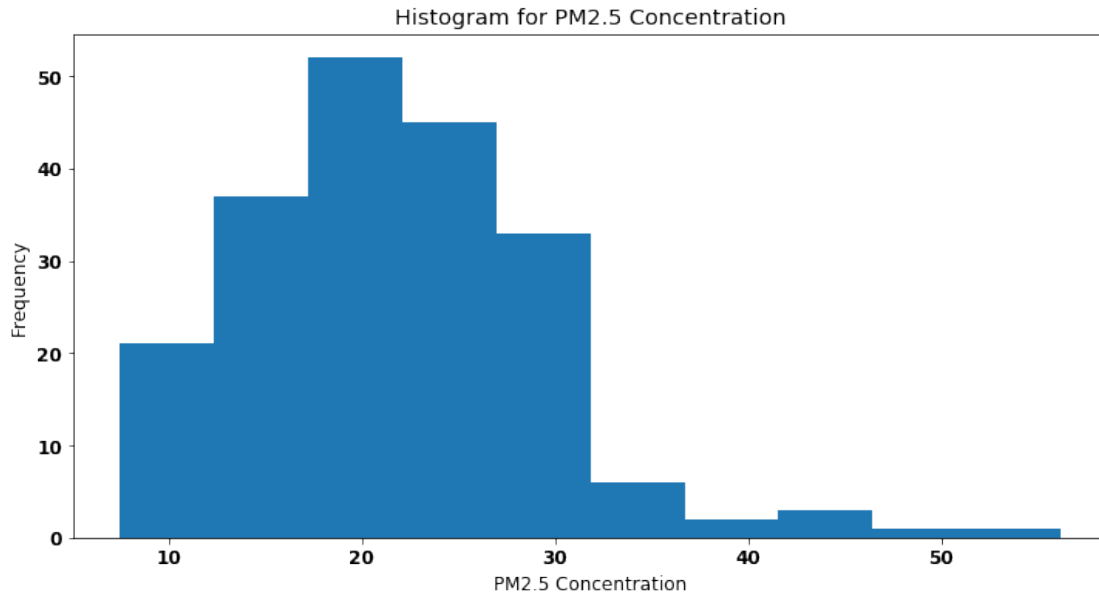


0.1.13 Daily Summaries

```
In [58]: print(daily_average_airquality_data_concentrations.shape) #shape of dataframe
          print(daily_average_airquality_data_concentrations.max()) #max values in dataframe
```

```
(201, 8)
s1_pm2_5    55.21
s1_pm10     69.79
s2_pm2_5    57.14
s2_pm10     67.31
Latitude     1.33
Longitude    34.37
pm2_5        56.18
pm10         68.55
dtype: float64
```

```
In [59]: plt.rc('figure', figsize=(12,6))
          plt.hist(daily_average_airquality_data_concentrations['pm2_5'], 10)
          plt.ylabel('Frequency')
          plt.xlabel('PM2.5 Concentration')
          plt.title('Histogram for PM2.5 Concentration')
          plt.show()
```



```
In [60]: fig = plt.figure(figsize=(14,6))


plt.rcParams.update({'font.size':14, 'font.weight':'bold'})

plt.plot(daily_average_airquality_data_concentrations.index,
         daily_average_airquality_data_concentrations['pm2_5'],
         color='olive', linestyle='solid', label='PM2.5 Average Values')

plt.plot(daily_maximum_airquality_data_concentrations.index,
         daily_maximum_airquality_data_concentrations['pm2_5'],
         color='chocolate', linestyle='solid', label='PM2.5 Maximum Values')

plt.plot(daily_minimum_airquality_data_concentrations.index,
         daily_minimum_airquality_data_concentrations['pm2_5'],
         color='darkcyan', linestyle='solid', label='PM2.5 Minimum Values')

chart_title = 'Daily PM2.5 Concentration for Location X in 2019'
plt.title(chart_title,fontsize=20)
plt.ylabel('PM2.5 Concentration', fontsize=20)
plt.xlabel('Date', fontsize=20)
plt.xticks(rotation=0)
plt.legend()
plt.show()
fig.savefig(os.path.join(base_dir,'daily_pm25_concentration_location_x_2019.png'))
```



```
In [61]: ind = np.arange(len(monthly_average_airquality_data_concentrations))
width = 0.36
#0.525
fig = plt.figure(figsize=(13,4))
ax = fig.add_subplot(111)
rects1 = ax.bar(ind- width/2, monthly_minimum_airquality_data_concentrations['pm2_5'])
rects3 = ax.bar(ind + width, monthly_average_airquality_data_concentrations['pm2_5'],
rects2 = ax.bar(ind+ width/4, monthly_maximum_airquality_data_concentrations['pm2_5'])

ax.set_ylabel('PM2.5 Concentration',fontSize=16)
ax.set_title(' Monthly PM2.5 Concentration For 2019', fontsize=16)
ax.set_xlabel('Date', fontsize=16)
ax.set_xticks(ind)
x_tick_labels = pd.to_datetime(monthly_average_airquality_data_concentrations.index)
final_labels= []
for x in x_tick_labels:
    final_labels.append(x.strftime('%b'))
ax.set_xticklabels(final_labels, fontsize=14)
ax.legend()

plt.show()
fig.savefig( 'monthly_pm25_concentration_bargraph.png')
```